

IMPERIAL



MENG INDIVIDUAL PROJECT

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Tools for Foundational Verification of Probabilistic Programs

Author:
Edwin Fernando

Supervisor:
Shing Hin Ho
Dr. Azalea Raad

Second Marker:
?

January 20, 2026

Contents

1	Introduction	3
1.1	Objectives	3
1.2	Challenges	3
1.3	Contributions	3
2	Background	4
2.1	Measure Theory	4
2.1.1	Pi-lambda theorem	5
2.1.2	Disintegration theorem	5
2.2	Semantics of Probabilistic Programs	5
2.2.1	Semantics of Let binding	5
2.2.2	Semantics of sample, observe and normalize	5
2.3	Lean and relevant recent formalisations	5
2.4	Separation Logic	6
2.5	Symbolic Execution tools	6
2.6	Iris and MoSeL	6
2.7	Related Work	7
3	Project Plan	8
4	Formalisation of Lilac	10
4.1	Measure spaces as a dependent type	10
4.2	Pi-lambda-theorem as an induction principle	10
4.3	Parametrisation of the "store" by the "heap" in this separation logic	10
5	Evaluation	11

Chapter 1

Introduction

Probabilistic programming languages (Hakaru, Stan, Gen, etc [1–3]) offer a principled way of creating statistical models complex probability distributions over multiple random variables, as code. The programmer (statistician) can then interact with the model, evaluating (sampling from) the generated distribution. This decouples the specification of a model from the inference algorithm used for sampling. Probabilistic programming is in fact also pervasive in computer science, though generally not cleanly abstracted into a DSL. Examples include Zero Knowledge Proof Systems (Blockchain) [4], Differential Privacy and Machine Learning (TensorFlow Probability [5]). Formal reasoning about programs in these critical applications may be highly beneficial.

The aim of this project is to extend the reach of foundational verification in this direction. The main contributions are:

1. **C1:** Implementing (formalising) Lilac [6] and BaSL [7] separation logics in the Lean theorem prover, with "tactical support", using Iris [8].
2. **C2:** (if time permits) Develop a theory of compositional symbolic execution (CSE) for probabilistic programs and implement using Soteira [9]. CSE can automate program verification.

C1 is one of the first (if not the first) foundationally verified (in Lean) implementation of a probabilistic separation logic (PSL). There is in fact a simultaneous project formalising the Bluebell PSL [10] for the purpose of verifying the core of a Zero Knowledge Proof system [4].

C2 would be a novel theoretical contribution. The main innovation here, would be either The separation logic chosen for this may be Lilac or BaSL, but also possibly Bluebell [11], to directly apply the work in an industrial formalisation project. However Bluebell supports both unary

with various subsystems But there are much wider applications in computer science, where arguments of correctness are probabilistic in nature. Examples expressively and formally verify these arguments if the implementation was carried out in a probabilistic programming language/DSL, using Probabilistic Separation Logics (PSL). This project is meant as a proof of concept, aiming to be the first implementation of a PSL. Specifically in the Lean proof assistant, we implement the lilac separation logic using Iris-lean (ported from rocq) to provide ergonomic tactical support. We choose Lean for the significant body of measure-theoretic results present in Mathlib, required to establish soundness of the logic.

Given time I would also like to pursue formalisation of BaSL or a theory for symbolic execution using lilac or BaSL for probabilistic programs.

on top of the Iris framework, in Lean. offer an intuitive way to reason about probabilistic programs at a suitable level of abstraction.

1.1 Objectives

1.2 Challenges

1.3 Contributions

Chapter 2

Background

Probabilistic programming is a paradigm where we are defining statistical models using code. It is a means of describing complex distributions over some space in a structured way. Probabilistic programming language (PPL) also come with an interpreter (probabilistic inference engine) which allows the programmer to sample a term from the distribution denoted by a program and indeed sample distributions while constructing new distributions.

More interestingly a Bayesian probabilistic programming language (BPPL) also allows one to "observe" an event and update their belief of the world (update the probability distribution they associate with some occurrence). Here we take the bayesian perspective underpinned by Baye's law which describes how we update a probability distribution based

TODO: buses and days and instantiation of Baye's law

A simple example taken from [12] is reproduced below to show how

2.1 Measure Theory

In order to understand the semantics of a probabilistic programming, a formal definition for probability is needed. Measure theory underpins probability theory and an overview is given here. We want to be able to talk about the probabilities of some event occurring, which is modelled by a **random variable**. In order to define a random variable we first need to define Measures and Measurable Space. A measurable space is a pair $(\Omega, \mathcal{F})$, where Ω is some Set, and $\mathcal{F} \subseteq \mathcal{P}(\Omega)$ is called the σ -algebra. We interpret each element of $\mathcal{F}$ to be an event (such as a coin flip landing heads). We would like the σ -algebra to closed under complements and countable unions. From the presentation so far it may seem like all singleton sets must be events but due to technical reasons (see [13] Ch. 2), this cannot be in certain cases where Σ is infinite.

A measure on a measurable space is a function $\mu : \mathcal{F} \rightarrow [0, \infty]$

This is essentially providing a probability density or probability mass function for the set of events $\mathcal{F}$. A similar but more general notion than what a measure gives us, is a random variable: a measurable function between two measurable spaces. I say "more general notion", since using the identity measurable function yields the measure we started with. A measurable function denoted $f : (\Omega_1, \mathcal{F}) \xrightarrow{m} (\Omega_2, \mathcal{G})$ $f : \Omega_1 \xrightarrow{m} \Omega_2$ is well defined when there is an ambient measurable space structure $(\Omega_1, \mathcal{F})$ and $(\Omega_2, \mathcal{G})$. f allows us to talk about the distribution of a particular attribute of in the event space. For example $\Omega_1 = [1, 6]^2$ may be the set of all events associated with playing 2 dice, while f may map to the sum of the two dice, $\Omega_2 = [1, 12]^2$. A measurable function f , transforms any measure μ on $(\Omega_1, \mathcal{F})$ to the pushforward measure μ' on $(\Omega_2, \mathcal{G})$. In order for this to always be well defined we require the following measurability to be satisfied by all measurable functions: $\forall G \in \mathcal{G}, f^{-1}(G) \in \mathcal{F}$.

TODO: Define lebesgue integration, since it's used in the Giry monad. TODO: Define s-finite measures and kernels

2.1.1 Pi-lambda theorem

2.1.2 Disintegration theorem

The disintegration theorem allows a more general definition conditional variables where conditioning on events with zero probability is well defined under certain conditions. This is required, since if conditioning is defined in a program logic, it must handle all edge cases

TODO: more details in final report (as I haven't studied this yet)

2.2 Semantics of Probabilistic Programs

Staton provided a compositional semantics to first-order Bayesian probabilistic programs in [12], which forms the foundation of program logics, validation of inference algorithms, etc. The scope of this project is entirely within first-order probabilistic programs (probability distributions over function spaces are not expressible). The programming language considered in Lilac uses a simpler semantics since it doesn't support bayesian updates. BaSL uses exactly this semantics. Note that supporting bayesian updates, involves the construct **observe** and **normalize** in the language. We will present here the more complete semantics. The key contribution of the paper is the introduction of **s-finite kernels**, motivated by validating a basic commutativity property that programmers expect. Specifically:

$$\begin{array}{ccc} \text{let } x = t \text{ in} & & \text{let } y = u \text{ in} \\ \text{let } y = u \text{ in} & = & \text{let } x = t \text{ in} \\ v & & v \end{array}$$

when neither x is free in u nor y is free in t . Another benefit of this semantics is a line-by-line compositionality, ie, the semantics of every subterm in a program is well-defined, not just whole programs. This is just another way of saying the semantics of open programs (those with free variables), as well as closed programs is well-defined, respectively. Validating proof rules in a logic is much better formulated in terms of such a compositional semantics. s-finite kernels denote open programs and s-finite measures denote closed programs.

2.2.1 Semantics of Let binding

Probabilistic programs are interpreted within the Giry monad [14]. Probability is an effect. This interpretation gives us the semantics of sequencing (let-binding) so we go through it below. Kernels also naturally show up in type of the monadic bind, which motivates it. The setting is the category of Measurable Spaces, **Meas**. All types in our language is interpreted as an object in this category. Syntactically:

$$\mathbb{A}, \mathbb{B} ::= \mathbb{R} \mid \mathcal{P}(\mathbb{A}) \mid 1 \mid \mathbb{A} \times \mathbb{B} \mid \dots$$

Objects in the category of the form (A, Σ_A) . Morphisms are measurable functions. The interesting type here is $\mathcal{P}(\mathbb{A})$. It's the type of probability distributions over $\mathbb{A}$. Actually not quite, since we allow any measure $\mu : \Sigma_A \rightarrow [0, \infty]$, rather than $\Sigma_A \rightarrow [0, 1]$. Now in order to be an object in the category, $\mathcal{P}(\mathbb{A})$, the set of all measures on $\mathbb{A}$. must itself form a measurable space. There is some construction to form a σ -algebra for $\mathcal{P}(\mathbb{A})$ satisfying all the requirements (we don't go into it here).

The bayesian probabilistic programming language (BPPL) is a standard ML like language with 3 special constructs, which are well-typed in the following cases:

We work in the categor Below I just state the **pure** and **bind** operations in the Giry monad, without proofs of the monad laws. To think about commutativity of let bindings we will have to look at the semantics of the sequencing operations in a probabi

2.2.2 Semantics of sample, observe and normalize

2.3 Lean and relevant recent formalisations

Foundational verification, consists of machine checked proofs of theorems, assuming only the axioms of mathematics. Here we are concerned with verification of software, so the theorems in question express properties of the software. The axioms of mathematics that we use is a dependent type

theory. In this setting of dependent type theory, checking the correctness of proof is the same thing as checking if a proof (program) is well typed, due to the Curry-Howard Isomorphism [CITE](#). Further the tool we use to for machine-checked proofs is an interactive theorem prover with support for proofs using "tactics". Tactical support is just metaprogramming, trying to encapsulate reasoning principles a person would use when writing pen-and-paper proofs. There are two proof assistants that were strongly considered for the purposes of this project: Lean and Rocq, former having proofs of required theorems in measure theory and the latter having the original implementation of the Iris framework. Lean was chosen since the front-end of the Iris proof interface (MoSeL, see [2.6](#)) had been fully ported to Lean very recently.

There has been an exciting amount of recent activity in formalisations related to this work, which have enabled the formalisation of Lilac (see, [??](#)) in Lean to be a realistic goal

A lot of background in Lean was required to undertake this project. The Lean development exercises the full expressivity of inductive types in Lean, for example and a thorough understanding of propositional vs definitional equality which is related to Proof irrelevance [Theorem proving for Lean4] and how that's implemented in Lean's type theory by special treatment of the lowest Universe [Mario's thesis]

Various representations of formal languages particularly the problem of how to represent variables in contexts, with strong guarantees from the types System. [Cite chlilala, HOAS, and lean-lang.org]

2.4 Separation Logic

O'Hearn's original paper. PSL, Lilac, BaSL

It is unclear to me why an axiomatic semantics would need to be provided in terms of measurable maps from the Hilbert cube. I wonder whether measure theory is not abstract enough and perhaps synthetic measure theory may hold some intuitions/inspiration or simplifications

2.5 Symbolic Execution tools

2.6 Iris and MoSeL

Iris [\[8\]](#) is a step-indexed modal separation logic partly designed as a framework for embedding other separation logics into, to take advantage of the Iris Proof Mode (IPM) [\[15\]](#). The Iris proof mode provides state of the art "tactical support", tactic based interactive theorem proving. More recently the front of the IPM was abstracted from Iris to any separation logic to create MoSeL [\[16\]](#). We aim to use MoSeL in this project for several reasons

1. Neither Lilac nor BaSL is step-indexed, so it is much cleaner to instantiate MoSeL rather than Iris
2. Making Lilac compatible with Lilac's interface (instantiating an algebraic structure called CMRA [\[8\]](#)) is an unsolved problem (future work section [\[6\]](#))
3. MoSeL was completely ported to Lean, shortly before the start of this project! [\[17\]](#).

MoSeL presents the interface of a "logic of Bunched Implications" which essentially contains all standard connectives of a first order logic and the separation logic connectives. The only addition MoSeL adds to this is a persistence modality (which can be instantiated in a trivial way in exchange for defeating any automation it might achieve?).

One of the key generalisations that MoSeL makes over Iris is parametricity over linear and affine logics. Iris is an affine logic meaning separating conjuncts in a premise can be used 0 or 1 time in proving a goal; they can be thrown away if desired. In a linear logic however every conjunct must be used exactly once, as in the first separation logic by O'Hearn and Pym [\[18\]](#). There is naturally quite an emphasis on describing whether a logic is affine or linear, since throwing away unwanted predicates when it is suitable to do so, is exactly the kind of thing a user wants automation for. In BaSL [\[7\]](#), there is a short discussion stating BaSL is part affine and part linear. Precisely characterising the constructs which are affine and linear or under what conditions should be explored for better tactical support.

MoSeL like IPM before it provides tactics mirroring the native Lean tactics, such as `iIntros` and `IAssumption`

In an ongoing project to port Iris from Rocq to Lean, MoSeL was completely ported to Lean very re

2.7 Related Work

TODO: List the quasi borel spaces formalisation and stuff

Chapter 3

Project Plan

I want to note that I have not before seen the interpretation of a typed calculus, taking both the type and typing context as parameters to give the denotation of a term. This definition of terms with they're typing context, actually massively simplifies the lean development previous iterations, involved a type inference algorithm, which could fail and the denotation function was partial.

I present 3 possible goals for the project here. M1 is the fallback which is mostly formalisation of already existing theory. M2 and M3 are more ambitious goals involving novel theoretical contributions. One of my original motivations for a formalisation project was to thoroughly learn and understand a topic in programming language theory, all the way up to the state of the art. Simultaneously I would be making a non-trivial contribution by formalising a program logic in a theorem prover and developing "tactical support" for it, showing a prototype for what is possible using foundational verification at the moment. On top of formally verifying the soundness of Lilac and BaSL, this could show how ergonomic or expressive a logic is. However this is still only a prototype because of the requirement to use the toy language (BPPL) This is only a "prototype"

The original scope of the project was M1 alone but t

In the case of M2 there is also a scope of industrial use of an implementation in Lean.

- Formalisation of Lilac with tactical support
- Lean interpreter for the language which is sound w.r.t to the denotational semantics
- The first case of Compositional Symbolic execution for probabilistic separation logic

CSE for the Bluebell separation logic with possible immediate application in industry for verifying correctness of Zero Knowledge Proof systems This could perhaps be a tactic, which actually invokes the symbolic execution to elaborate into a sequence of proof rules?

This is possibly not so hard


CSE for Lilac or BaSL. this is perhaps out of scope of project. But it connects to the first part of the project formalising Lilac/BaSL which is novel for being the first mechanisation of a separation logic for continuous probabilistic programs.

Put simply my goal is to reach one of the orange boxes. The "full proof of concept" is fallback goal. The major contribution here would be the formalisation of BaSL in Lean, which has several unknown elements of pure math proofs to it, when compared to to the Lilac formalisation. To name a few 1) Pi-lambda theorem (elaborate) 2) Disintegration theorem

The goal "industrial application" requires the development of a theory of compositional symbolic execution using Bluebell as the assertion language. So even if the Lean formalisations I am dependent on to implement a tool, using the theory is not available in time, I can still make a significant theoretical contribution here.

is dependent on completion of work planned by others, by the time I reach that stage of my project.

Figure 3.1: Example of a parametric plot $(\sin(x), \cos(x), x)$



Chapter 4

Formalisation of Lilac

Here I present the current progress in the formalisation of Lilac, and how I went about it. I divide it into 2 parts Initial work which was mostly measure theoretic, and recent work which is more structured. As in any project its best to get to the most uncertain parts of it first, to minimise its risk. I think I have somewhat intuitively followed this principle in my formalisation so far, also meaning its full of `sorry`s. However the shape of the formalisation of Lilac is reasonably clear and I was able to identify some noteworthy problems which show up in the formalisation.

As a brief reflection on the difference in formalisation on the maths and computer science sides, in my limited experience, I do think in this project (computer science) I do need to think of multiple things at once and papers generally disseminate dependencies more clearly without depending as heavily on the reader to morph and modify lemmas as required. However the objects in Lilac are so unconventional from the measure theory in `mathlib`, that it requires inventing some things from scratch and quite challenging. This is ideal for my goals of understanding how to use a dependent type theory for formalisation. As such [4.2](#) talking about formalisation specific issues. But [4.3](#) talks about something more conceptual in Lilac! (My pivot towards the aim of a more theoretical contribution is that I feel I have gotten enough Knowledge of Lean from formalising Lilac)

4.1 Measure spaces as a dependent type

4.2 Pi-lambda-theorem as an induction principle

The $\pi - \lambda$ -theorem is tool in measure theory for showing equality of two measures. It can be thought of as an induction principle, (indeed this is how it is defined in Lean) on the structure of a measure. Specifically, the definition

4.3 Parametrisation of the "store" by the "heap" in this separation logic

- The correct type for measurable spaces because type casting issue, it needed to be a dependent type. - The inductive family definition of a sigma algebra, and the pi lambda theorem being an induction problem complicates thing when we can have multiple induction principles and we really want a specific induction principle! - Correctly instantiating BI, modelling random variables, "intrinsically type calculus"? - Talk about Lilac, give the intuition that this is the actual line by line semantics

Chapter 5

Evaluation

Bibliography

- [1] Praveen Narayanan, Jacques Carette, Wren Romano, Chung-chieh Shan, and Robert Zinkov. Probabilistic inference by program transformation in hakaru (system description). In *International Symposium on Functional and Logic Programming - 13th International Symposium, FLOPS 2016, Kochi, Japan, March 4-6, 2016, Proceedings*, pages 62–79. Springer, 2016. doi: 10.1007/978-3-319-29604-3_5. URL http://dx.doi.org/10.1007/978-3-319-29604-3_5.
- [2] Bob Carpenter, Andrew Gelman, Matthew D. Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. Stan: A probabilistic programming language. *Journal of Statistical Software*, 76(1):1–32, 2017. doi: 10.18637/jss.v076.i01. URL <https://www.jstatsoft.org/index.php/jss/article/view/v076i01>.
- [3] Marco F. Cusumano-Towner, Feras A. Saad, Alexander K. Lew, and Vikash K. Mansinghka. Gen: a general-purpose probabilistic programming system with programmable inference. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2019*, page 221–236, New York, NY, USA, 2019. Association for Computing Machinery. ISBN 9781450367127. doi: 10.1145/3314221.3314642. URL <https://doi.org/10.1145/3314221.3314642>.
- [4] Verified-zkEVM. Arklib: Formally verified arguments of knowledge in lean, 2025. URL <https://github.com/Verified-zkEVM/ArkLib>. Accessed: 2025-12-15.
- [5] Joshua V. Dillon, Ian Langmore, Dustin Tran, Eugene Brevdo, Srinivas Vasudevan, Dave Moore, Brian Patton, Alex Alemi, Matt Hoffman, and Rif A. Saurous. Tensorflow distributions, 2017. URL <https://arxiv.org/abs/1711.10604>.
- [6] John M. Li, Amal Ahmed, and Steven Holtzen. Lilac: A modal separation logic for conditional probability. *Proc. ACM Program. Lang.*, 7(PLDI), June 2023. doi: 10.1145/3591226. URL <https://doi.org/10.1145/3591226>.
- [7] Shing Hin Ho, Nicolas Wu, and Azalea Raad. Bayesian separation logic, 2025. URL <https://arxiv.org/abs/2507.15530>.
- [8] Ralf Jung, Robbert Krebbers, Jacques-henri Jourdan, Aleš Bizjak, Lars Birkedal, and Derek Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation logic. *Journal of Functional Programming*, 28:e20, 2018. doi: 10.1017/S0956796818000151.
- [9] Sacha-Élie Ayoun, Opale Sjöstedt, and Azalea Raad. Soteria: Efficient symbolic execution as a functional library, 2025. URL <https://arxiv.org/abs/2511.08729>.
- [10] Verified-zkEVM. iris-lean. <https://github.com/Verified-zkEVM/iris-lean>, 2025. Fork of the Iris Lean repository, instantiating a probabilistic separation logic, Bluebell.
- [11] Jialu Bao, Emanuele D’Osualdo, and Azadeh Farzan. Bluebell: An alliance of relational lifting and independence for probabilistic reasoning. *Proc. ACM Program. Lang.*, 9(POPL), January 2025. doi: 10.1145/3704894. URL <https://doi.org/10.1145/3704894>.
- [12] Sam Staton. Commutative semantics for probabilistic programming. In *Programming Languages and Systems: 26th European Symposium on Programming, ESOP 2017, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2017, Uppsala, Sweden, April 22–29, 2017, Proceedings*, page 855–879, Berlin, Heidelberg, 2017. Springer-Verlag. ISBN 978-3-662-54433-4. doi: 10.1007/978-3-662-54434-1_32. URL https://doi.org/10.1007/978-3-662-54434-1_32.

- [13] Sheldon Axler. *Measures*, pages 13–72. Springer International Publishing, Cham, 2020. ISBN 978-3-030-33143-6. doi: 10.1007/978-3-030-33143-6_2. URL https://doi.org/10.1007/978-3-030-33143-6_2.
- [14] Michèle Giry. A categorical approach to probability theory. In B. Banaschewski, editor, *Categorical Aspects of Topology and Analysis*, pages 68–85, Berlin, Heidelberg, 1982. Springer Berlin Heidelberg. ISBN 978-3-540-39041-1.
- [15] Robbert Krebbers, Amin Timany, and Lars Birkedal. Interactive proofs in higher-order concurrent separation logic. *SIGPLAN Not.*, 52(1):205–217, January 2017. ISSN 0362-1340. doi: 10.1145/3093333.3009855. URL <https://doi.org/10.1145/3093333.3009855>.
- [16] Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser, Amin Timany, Arthur Charguéraud, and Derek Dreyer. Mosel: a general, extensible modal framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.*, 2(ICFP), July 2018. doi: 10.1145/3236772. URL <https://doi.org/10.1145/3236772>.
- [17] leanprover-community. iris-lean: Lean 4 port of iris, a higher-order concurrent separation logic framework. <https://github.com/leanprover-community/iris-lean/>, 2026. Accessed: 18th Jan 2026.
- [18] Peter W. O’Hearn and David J. Pym. The logic of bunched implications. *Bulletin of Symbolic Logic*, 5(2):215–244, 1999. doi: 10.2307/421090.